Image processing, analysis and and machine vision (3. ed.).

Book · January 2008

Source: DBLP

Vaclav Hlavac
Czech Technical University in Prague

Roger David Boyle
Aberystwyth University

Image Processing, Analysis, and Machine Vision

Second Edition

Milan Sonka   Vaclav Hlavac   Roger Boyle

Contents

List of algorithms

List of symbols and abbreviations

$\arg(x, y)$	angle (in radians) from x axis to the point (x, y)
$\mathrm{argmax}_i(\mathrm{expr}(i))$	the value of i that causes $\mathrm{expr}(i)$ to be maximal
$\mathrm{argmin}_i(\mathrm{expr}(i))$	the value of i that causes $\mathrm{expr}(i)$ to be minimal
div	integer division
mod	remainder after integer division
$\mathrm{round}(x)$	largest integer which is not bigger than $x + 0.5$
$\emptyset$	empty set
A^c	complement of the set A
$A \subset B$, $B \supset A$	set A is included in set B
$A \cap B$	intersection between sets A and B
$A \cup B$	union of sets A and B
$A \mid B$	difference between sets A and B
$\mathbf{A}$	(uppercase bold) matrices
$\mathbf{x}$	(lowercase bold) vectors
$\lvert\, \mathbf{x} \,\rvert$	magnitude (or modulus) of the vector $\mathbf{x}$
$\mathbf{x} \cdot \mathbf{y}$	scalar product between vectors $\mathbf{x}$ and $\mathbf{y}$
$\tilde{x}$	estimate of the value x
$\lvert\, x \,\rvert$	absolute value of a scalar
$\Delta\, x$	small finite interval of x, difference
$\partial f / \partial x$	partial derivative of the function f with respect to x
$\nabla\, \mathbf{f}$, $\mathrm{grad}\, \mathbf{f}$	gradient of $\mathbf{f}$
$\nabla^2\, \mathbf{f}$	Laplace operator applied to $\mathbf{f}$
$f * g$	convolution between functions f and g
D_E	Euclidean distance (see Section 2.3.1)
D_4	city block distance (see Section 2.3.1)
D_8	chessboard distance (see Section 2.3.1)
F	complex conjugate of the complex function F
$\mathrm{rank}(A)$	Rank of a matrix A
T^*	transformation dual to transformation T
$\delta(x)$	Dirac function
$\mathcal{E}$	mean value operator
$\mathcal{L}$	linear operator

$\mathcal{O}$	origin of the coordinate system
$\#$	number of (e.g., pixels)
$\breve{B}$	point set symmetrical to point set B
$\oplus$	morphological dilation
$\ominus$	morphological erosion
$\circ$	morphological opening
$\bullet$	morphological closing
$\otimes$	morphological hit-or-miss transformation
$\oslash$	morphological thinning
$\odot$	morphological thickening

1D	one dimension(al)
2D	two dimension(al)
3D	three dimension(al)
AI	artificial intelligence
ASM	active shape model
B-rep	boundary representation
CAD	computer-aided design
CCD	charge-coupled device
CSG	constructive solid geometry
CT	computer tomography
dof	degrees of freedom
ECG	electro-cardiogram
EEG	electro-encephalogram
FFT	fast Fourier transform
FOE	focus of expansion
GA	genetic algorithm
HMM	hidden Markov model
IHS	intensity, hue, saturation
JPEG	Joint Photographic Experts Group
MR	magnetic resonance
MRI	magnetic resonance imaging
OCR	optical character recognition
OS	order statistics
PDM	point distribution model
PET	positron emission tomography
PMF	Pollard-Mayhew-Frisby (correspondence algorithm)
RGB	red, green, blue
SNR	signal-to-noise ratio
SVD	singular value decomposition
TV	television

Preface

Image processing, analysis and machine vision represent an exciting and dynamic part of cognitive and computer science. Following an explosion of interest during the 1970s, the 1980s and 1990s were characterized by the maturing of the field and the significant growth of active applications; remote sensing, technical diagnostics, autonomous vehicle guidance, medical imaging (2D and 3D) and automatic surveillance are the most rapidly developing areas. This progress can be seen in an increasing number of software and hardware products on the market, as well as in a number of digital image processing and machine vision courses offered at universities worldwide.

There are many texts available in the areas we cover—most (indeed, all of which we know) are referenced somewhere in this book. The subject suffers, however, from a shortage of texts which are 'complete' in the sense that they are accessible to the novice, of use to the educated, and up to date. Here we present the second edition of a text first published in 1993 in which we hope to include many of the very rapid developments that have taken and are still taking place, which quickly age some of the very good textbooks produced over the last two decades or so. The target audience is the undergraduate with negligible experience in the area through to the Master's and research student seeking an advanced springboard in a particular topic. Every section of this text has been updated since the first version (particularly with respect to references); additionally, wholly new sections are presented on: compression via JPEG and MPEG; fractals; fuzzy logic recognition; hidden Markov models; Kalman filters; point distribution models; three-dimensional vision; watershed segmentation; wavelets; and an entire chapter devoted to case studies. Additionally, each chapter now includes a concise Summary section. To help the reader to acquire practical understanding, newly added Exercise sections accompany each chapter; these are in the form of short-answer questions and problems of varying difficulty, frequently requiring practical usage of computer tools and/or development of application programs.

This book reflects the authors' experience in teaching one- and two-semester undergraduate and graduate courses in Digital Image Processing, Digital Image Analysis, Machine Vision, Pattern Recognition, and Intelligent Robotics at their respective institutions. We hope that this combined experience will give a thorough grounding to the beginner and provide material that is advanced enough to allow the more mature student to understand fully the relevant areas of the subject. We acknowledge that in a very short time the more active areas will have moved beyond this text.

This book could have been arranged in many ways. It begins with low-level processing and works its way up to higher levels of image interpretation; the authors have chosen this framework because they believe that image understanding originates from a common database of information. The book is formally divided into 16 chapters, beginning with low-level processing and working toward higher-level image representation, although this structure will be less apparent after Chapter 10, when we present transforms, compression, morphology, texture, and motion analysis which are very useful but often special-purpose approaches that may not always be included in the processing chain. The final chapter presents four live research projects which illustrate in practical use much of what has gone before.

Decimal section numbering is used, and equations and figures are numbered within each chapter. Each chapter is accompanied by an extensive list of references and exercises. A selection of algorithms is summarized formally in a manner that should aid implementation—not all the algorithms discussed are presented in this way (this might have doubled the length of the book); we have chosen what we regard as the key, or most useful or illustrative, examples for this treatment.

Each chapter presents material from an introductory level through to an overview of current work; as such, it is unlikely that the beginner will, at the first reading, expect to absorb all of a given topic. Often it has been necessary to make reference to material in later chapters and sections, but when this is done an understanding of material in hand will not depend on an understanding of that which comes later. It is expected that the more advanced student will use the book as a reference text and signpost to current activity in the field—we believe at the time of going to press that the reference list is full in its indication of current directions, but record here our apologies to any work we have overlooked. The serious reader will note that many references are very recent, and should be aware that before long more relevant work will have been published that is not listed here.

This is a long book and therefore contains material sufficient for much more than one course. Clearly, there are many ways of using it, but for guidance we suggest an ordering that would generate four distinct modules:

Digital Image Processing, an undergraduate course.

Digital Image Analysis, an undergraduate/graduate course, for which Digital Image Processing may be regarded as prerequisite.

Computer Vision I, an undergraduate/graduate course, for which Digital Image Processing may be regarded as prerequisite.

Computer Vision II, a graduate course, for which Computer Vision I may be regarded as prerequisite.

The important parts of a course, and necessary prerequisites, will naturally be specified locally; a suggestion for partitioning the contents follows this Preface.

Assignments should wherever possible make use of existing software; it is our experience that courses of this nature should not be seen as 'programming courses', but it is the case that the more direct practical experience the students have of the material discussed, the better is their understanding. Since the first edition was published, an explosion of World Wide Web-based material has been made available, permitting many of the exercises we present to be conducted without the necessity of implementing from scratch—we do not present explicit pointers to Web material, since they evolve so quickly; however, pointers to specific support materials for this book and others may be located via the publisher, *http://www.brookscole.com.*

The book has been prepared using the LaTeX text processing system. Its completion would have been impossible without extensive usage of the Internet computer network and electronic mail. We should like to acknowledge the University of Iowa, the Czech Technical University, and the School of Computer Studies at Leeds University for providing the environment in which this book was prepared.

Milan Sonka was a faculty member of the Department of Control Engineering, Faculty of Electrical Engineering, Czech Technical University, Prague, Czech Republic for ten years, and is now an Associate Professor at the Department of Electrical and Computer Engineering, the University of Iowa, Iowa City, Iowa, USA. His research interests include medical image analysis, knowledge-based image analysis, and machine vision. Vaclav Hlavac is an Associate Professor at the Department of Control Engineering, Czech Technical University, Prague. His research interests are knowledge-based image analysis and 3D model-based vision. Roger Boyle is a Senior Lecturer in Artificial Intelligence in the School of Computer Studies at the University of Leeds, England, where his research interests are in low-level vision and pattern recognition. The first two authors have worked together for some years, and have been co-operating with the third since 1991.

The authors have spent many hours in discussions with their teachers, colleagues, and students, from which many improvements to early drafts of this text resulted. Particular thanks are due to Tomáš Pajdla, Petr Kodl, Radim Šára at the Czech Technical University; Steve Collins at the University of Iowa; Jussi Parkkinen at the University of Lappeenranta; Guido Prause at the University of Bremen; David Hogg at the University of Leeds; and many others whose omission from this list does not diminish the value of their contribution. The continuous support and encouragement we received from our wives and families, while inexplicable, was essential to us throughout this project—once again, we promise that our next book will not be written outside standard office hours or during holidays (but this time we mean it).

All authors have contributed throughout—the ordering on the cover corresponds to the weight of individual contribution. Any errors of fact are the joint responsibility of all, while any errors of typography are the responsibility of Roger Boyle. Jointly, they will be glad to incorporate any corrections into future editions.

Milan Sonka (milan-sonka@uiowa.edu)
The University of Iowa, Iowa City, Iowa, USA

Vaclav Hlavac (hlavac@vision.felk.cvut.cz)
Czech Technical University, Prague, Czech Republic

Roger Boyle (roger@scs.leeds.ac.uk)
University of Leeds, Leeds, England

Course contents

In this section, one possible ordering of the material covered in the four courses proposed in the Preface is given. This coverage should not be considered the only possibility—on the contrary, the possibilities for organizing Image Processing and Analysis courses are practically endless. Therefore, what follows should only be regarded as suggestions, and instructors should tailor course content to fit the already acquired knowledge, abilities, and needs of the students enrolled.

Digital Image Processing. An undergraduate course.

1 Introduction

2 The digitized image and its properties

3 Data structures for image analysis

4 Image pre-processing (excluding 4.3.6–4.3.9, 4.4.3, limited coverage of 4.3.4, 4.3.5)

5 Segmentation

- 5.1 Thresholding (excluding 5.1.3, 5.1.4)
- 5.2 Edge-based segmentation (excluding 5.2.8, limited coverage of 5.2.4, 5.2.5)
- 5.3 Region growing segmentation (excluding 5.3.4)
- 5.4 Matching

12 Linear discrete image transforms

13 Image data compression

16 Case studies (selected topics)

Digital Image Analysis. An undergraduate/graduate course, for which Digital Image Processing may be regarded as prerequisite. Sections that were covered in the Digital Image Processing class and re-appear are intended to be discussed at more depth than is possible in the introductory course.

1 Introduction (brief review)

2 The digitized image and its properties (brief review)

5 Segmentation

- 5.1.3 Multi-spectral thresholding
- 5.1.4 Thresholding in hierarchical data structures

- 5.2.4 Edge following as graph searching
- 5.2.5 Edge following as dynamic programming
- 5.3.4 Watershed segmentation

6 Shape representation and description (excluding 6.2.7, 6.3.4–6.3.6, 6.4)

7 Object recognition

- 7.1 Knowledge representation
- 7.2 Statistical pattern recognition
- 7.3 Neural networks
- 7.4 Syntactic pattern recognition

11 Mathematical morphology

14 Texture

16 Case studies (selected topics)

Computer Vision I. An undergraduate/graduate course, for which Digital Image Processing may be regarded as prerequisite.

1 Introduction (brief review)

2 The digitized image and its properties (brief review)

4 Image pre-processing

- 4.3.3 Zero-crossings of the second derivative
- 4.3.4 Scale in image processing
- 4.3.5 Canny edge detection
- 4.3.6 Parametric edge models
- 4.3.7 Edges in multi-spectral images
- 4.3.8 Other local pre-processing operators
- 4.3.9 Adaptive neighborhood pre-processing

6 Shape representation and description

7 Object recognition

8 Image understanding

16 Case studies (selected topics)

Computer Vision II. A graduate course, for which Computer Vision I may be regarded as prerequisite.

5 Segmentation

- 5.2.4 Edge following as graph searching
- 5.2.5 Edge following as dynamic programming
- 5.5 Advanced border and surface detection approaches

9 3D Vision, geometry and radiometry

10 Use of 3D vision

15 Motion analysis

Practical 3D vision projects

Index

Bold text refers to major or defining entries.